

GALETTE: a Lightweight XDP Dataplane on your Raspberry Pi

Kyle A. Simpson, Chris Williamson, Douglas J. Paul, Dimitrios P. Pazaros

✉ kylesimpson1@acm.org


🌐 FelixMcFelix <https://mcfelix.me>

13th June, 2023

University of Glasgow



Fast, cheap, and secure IoT Defence – pick 3?



- Single-board compute like RPi's are small, capable, affordable! **Cheap!**
 - See also: NUCs (££), Jetsons (£££).
 - *Linux-based:* Easy(/ier) to target and write for. **We also get kernel network stack advancements.**
 - Different CPU architectures.
- Sensor networks have low data rates; a good fit.
- Project goals:
 - **Fast!** Low-latency, quickly reconfigurable.
 - **Secure!** efficient NFV code gen from *memory-safe languages*.

GALETTE's Research Objectives

1. What specialisations does XDP Function Chaining need to best suit SBCs?
 - 'Acceptably' low-latency packet-processing, without pushing CPU/power draw too high?
2. How do we make eBPF + native compile from memory-safe systems languages easy? And portable across 'native'?
 - One Rust program per NF \implies compiles to eBPF + ***\$PLATFORM***.
 - Simple, dynamic chain format.
 - Fast reconfiguration.
3. How much better is it [power, perf, lat]?
 - With/without polling.

Background

Limits of existing SFC


- ‘Best’ low latency processing (DPDK) is **expensive** – CPU and power.
 - ...IFF you have HW support (NUCs)
- No powerful hardware offloads or acceleration.
 - FPGA hats/daughterboards ‘**off-path**’
- Devices physically vulnerable, **no ECC memory**.
- ...So, how to reconcile with cheap & portable SBCs?

eBPF: What and Why?

- Simple register machine VM (user-written) code, derived from BPF.
- Modern use – Kernel hooks, perf instrumentation, debugging
- JIT compiled
- Kernel-verified
 - Bounds-checked pointer accesses
 - Program size limited, no unbounded loops
 - Syscalls (*eBPF helpers*) exposed based on hook point



Network stack improvements: XDP



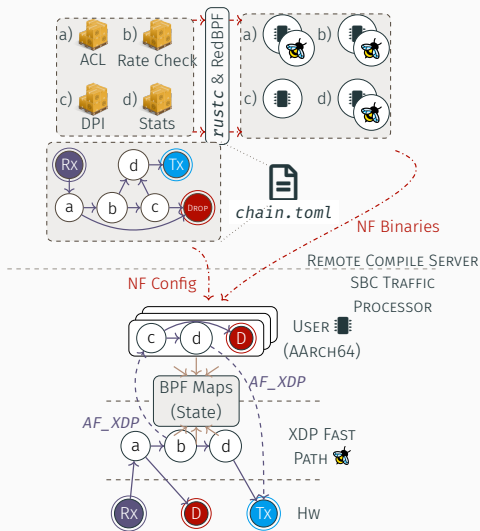
- eBPF hook attached to **packet ingress**
- Variations on hook
∈ {Offload, Driver, Generic}
 - Perf degrades gracefully according to driver support
- Hook can modify & inspect packets before forwarding to Linux stack, sending **straight to (another) NIC**, or drop.
- Since 2019: **AF_XDP** stack bypass!

Q1: Specialising ***AF_XDP*** SFC to SBCs

Concrete design differences

- **Problem:** Mismatch of HW queues to physical cores:
 - **Soln:** load balance or place high-latency NFs in userland.
 - ...also, don't pass packets back to k-space.
- **Problem:** XDP hooks only on ingress (*for now*):
 - **Soln:** load balance or place high-latency NFs in userland?
 - Write an individual NF *once*, compile for both envs, and replicate NFs as needed.

Design: Bird's eye view



- Two-tier approach—XDP & User.
- Composable NFs – graph structure.
- Critical or high performance NFs go into XDP:
 - Early results – low latency for most packets.
- Rare 'slow-path' still kernel bypass:
 - Expensive & proprietary code.
 - Only for candidate attack traffic.
- Reconfigurable, dynamic.

How does this differ from other frameworks?

In Security? SafeBricks¹, AuditBox² or similar.

- ...No SGX support in devices of interest.

In eBPF/XDP space? Polycube³!

- Built around datacentres – we often have just one HW queue for a NIC.

¹Poddar *et al.*, 'SafeBricks: Shielding Network Functions in the Cloud'.

²Liu *et al.*, 'Don't Yank My Chain: Auditable NF Service Chaining'.

³Miano *et al.*, 'A Framework for eBPF-Based Network Functions in an Era of Microservices'.

How do we upcall to userland?

- **Problem:** Can send packet over *AF_XDP*, but *no context on what the next (callee) NF is*.
 - Polycube's solution inadequate: one discrete userland component per *cube*.
- **Soln:** Adjust headroom of packets, write in ID and action of caller.
- ...might be a *memcpy*, but ideally only paid on packets who need it.

Q2: Easy Joint-Compile (eBPF +
Native) from Rust 🦀

Skeleton details

- Consistent NF API for both XDP/userland.
- Rust compiler should be able to enforce...
 - `#![forbid(unsafe_code)]` (or similar cargo tooling) on NF module crates,
 - all NF branches specified.
- All compilation on external server.
 - SBC too constrained.
 - If compile-server is TEE-equipped, can attest compiler/code etc. following SotA!

```
#![no_std]
pub enum Action {
    Left,
    Right,
    Up,
    Down,
}

pub fn packet(bytes: impl Packet) -> Action {
    let addr_lsb_idx = 14 +
    match pkt.slice_from(12, 2) {
        Some(&[0x08, 0x00]) => 19, //v4
        Some(&[0x86, 0xDD]) => 39, //v6
        _ => {return Action::Left},
    };

    match pkt.slice_from(addr_lsb_idx, 1)
        .map(|v| v[0] % 2) {
        Some(0) => Action::Left,
        Some(1) => Action::Right,
        Some(2) => Action::Up,
        Some(3) => Action::Down,
        _ => unreachable!(),
    }
}
```

mod.rs: Load balance on dest addr

< In lieu of a demo... >

Q3: Performance

?? Baselines

?? What Machines

?? What NFs.

?? WHY

?? Results?

content...

?? better at these things

If you want more detailed data, please check out our paper

Takeaways:

Cheap NFs: SBCs for packet processing.

Low-latency and fast: XDP path for majority of traffic, early & cheap anomaly checks, power savings.

Secure: PUFs for device, server, and function chain attestation.

Easy to write: *native and XDP* portable NFs in Rust.

Questions?

✉ kylesimpson1@acm.org

🌐 FelixMcFelix 🌐 <https://mcfelix.me>